


Support Vector Machines




Agenda

- Hard-margin SVM
 - theory (brief)
 - implementation
- Soft-margin SVM
 - theory (brief)
 - implementation
- Closing remarks

https://tiny.cc/mlt_workshop


- Colabs_Slides_Day_3
 - SVM.ipynb
 - Make a copy
- Code along

From Perceptrons to SVM




For linearly separable data with γ margin:

From Perceptrons to SVM



For linearly separable data with γ margin:
(1) Infinite number of valid linear classifiers.


From Perceptrons to SVM



For linearly separable data with γ margin:

- (1) Infinite number of valid linear classifiers.
- (2) Perceptron returns a valid linear classifier.


From Perceptrons to SVM



For linearly separable data with γ margin:

- (1) Infinite number of valid linear classifiers.
- (2) Perceptron returns a valid linear classifier.
- (3) Is it the "best"?


From Perceptrons to SVM





For linearly separable data with γ margin:

- (1) Infinite number of valid linear classifiers.
- (2) Perceptron returns a valid linear classifier.
- (3) Is it the "best"?
- (4) What is a good notion of "best"?

Margin





Margin




Small margin
Doesn't generalize well

Margin




Small margin
Doesn't generalize well




Large margin
Better generalization

Computing the Margin



For any linear classifier represented by \mathbf{w} :


Computing the Margin



For any linear classifier represented by \mathbf{w} :

(1) Find the point closest to it $\rightarrow \mathbf{x}^*$


Computing the Margin



For any linear classifier represented by \mathbf{w} :

- (1) Find the point closest to it $\rightarrow \mathbf{x}^*$
- (2) Scale \mathbf{w} such that \mathbf{x}^* lies on $\mathbf{w}^T \mathbf{x} = 1$


Computing the Margin



For any linear classifier represented by \mathbf{w} :

- (1) Find the point closest to it $\rightarrow \mathbf{x}^*$
- (2) Scale \mathbf{w} such that \mathbf{x}^* lies on $\mathbf{w}^T \mathbf{x} = 1$
- (3) Distance of \mathbf{x}^* from the line is $\frac{1}{\|\mathbf{w}\|}$


Computing the Margin



For any linear classifier represented by \mathbf{w} :


- (1) Find the point closest to it $\rightarrow \mathbf{x}^*$
- (2) Scale \mathbf{w} such that \mathbf{x}^* lies on $\mathbf{w}^T \mathbf{x} = 1$
- (3) Distance of \mathbf{x}^* from the line is $\frac{1}{\|\mathbf{w}\|}$
- (4) This is the (geometric) margin for this linear classifier.

Beyond the "margin"



$$(\mathbf{w}^T \mathbf{x}_i) y_i \geq 1, \quad 1 \leq i \leq n$$

Max-Margin Classifier




$$\max_{\mathbf{w}} \quad \frac{1}{\|\mathbf{w}\|}$$

sub. to

$$(\mathbf{w}^T \mathbf{x}_i) y_i \geq 1, \quad 1 \leq i \leq n$$

Max-Margin Classifier



$$\min_{\mathbf{w}} \frac{\|\mathbf{w}\|^2}{2}$$

sub. to

$$(\mathbf{w}^T \mathbf{x}_i) y_i \geq 1, \quad 1 \leq i \leq n$$

Primal and Dual

$$\begin{aligned} \min_{\mathbf{w}} \quad & \frac{||\mathbf{w}||^2}{2} \\ \text{sub. to} \quad & \end{aligned} \quad \equiv \quad \begin{aligned} \min_{\mathbf{w}} \max_{\boldsymbol{\alpha} \geq 0} \quad & \frac{||\mathbf{w}||^2}{2} + \sum_{i=1}^n \alpha_i \left[1 - (\mathbf{w}^T \mathbf{x}_i) y_i \right] \end{aligned} \quad \equiv \quad \begin{aligned} \max_{\boldsymbol{\alpha} \geq 0} \min_{\mathbf{w}} \quad & \frac{||\mathbf{w}||^2}{2} + \sum_{i=1}^n \alpha_i \left[1 - (\mathbf{w}^T \mathbf{x}_i) y_i \right] \end{aligned}$$
$$(\mathbf{w}^T \mathbf{x}_i) y_i \geq 1, \quad 1 \leq i \leq n$$
$$\boldsymbol{\alpha} = \begin{bmatrix} \alpha_1 \\ \vdots \\ \alpha_n \end{bmatrix}$$

Formulating the Dual

$$\max_{\boldsymbol{\alpha} \geq 0} \min_{\mathbf{w}} \frac{\|\mathbf{w}\|^2}{2} + \sum_{i=1}^n \alpha_i [1 - (\mathbf{w}^T \mathbf{x}_i) y_i]$$

$$\mathbf{w} = \sum_{i=1}^n \alpha_i \mathbf{x}_i y_i$$

Formulating the Dual

$$\max_{\boldsymbol{\alpha} \geq 0} \min_{\mathbf{w}} \frac{\|\mathbf{w}\|^2}{2} + \sum_{i=1}^n \alpha_i [1 - (\mathbf{w}^T \mathbf{x}_i) y_i]$$

$$\mathbf{w} = \sum_{i=1}^n \alpha_i \mathbf{x}_i y_i$$

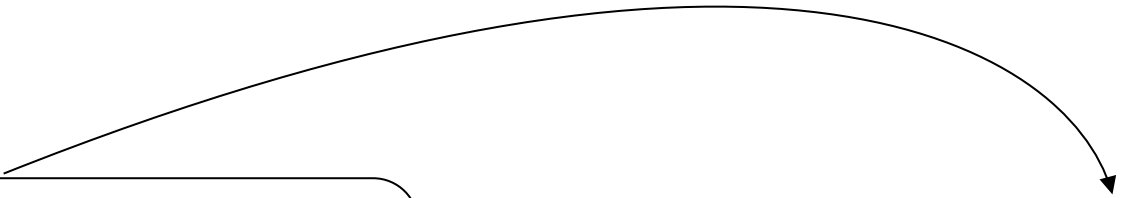
$$\mathbf{X} = \begin{bmatrix} | & & | \\ \mathbf{x}_1 & \cdots & \mathbf{x}_n \\ | & & | \end{bmatrix}$$

$$\boldsymbol{\alpha} = \begin{bmatrix} \alpha_1 \\ \vdots \\ \alpha_n \end{bmatrix}$$

$$\mathbf{Y} = \begin{bmatrix} y_1 & & 0 \\ & \ddots & \\ 0 & & y_n \end{bmatrix}$$

$$\mathbf{XY}\boldsymbol{\alpha} = \begin{bmatrix} | & & | \\ y_1 \mathbf{x}_1 & \cdots & y_n \mathbf{x}_n \\ | & & | \end{bmatrix} \begin{bmatrix} \alpha_1 \\ \vdots \\ \alpha_n \end{bmatrix} = \sum_{i=1}^n \alpha_i (y_i \mathbf{x}_i) = \mathbf{w}$$

Formulating the Dual

$$\max_{\boldsymbol{\alpha} \geq 0} \min_{\mathbf{w}} \frac{\|\mathbf{w}\|^2}{2} + \sum_{i=1}^n \alpha_i [1 - (\mathbf{w}^T \mathbf{x}_i) y_i]$$


$$\mathbf{w} = \mathbf{X}\mathbf{Y}\boldsymbol{\alpha}$$

$d \times n$

$$\mathbf{X} = \begin{bmatrix} | & & | \\ \mathbf{x}_1 & \cdots & \mathbf{x}_n \\ | & & | \end{bmatrix}$$

$$\boldsymbol{\alpha} = \begin{bmatrix} \alpha_1 \\ \vdots \\ \alpha_n \end{bmatrix}$$

$n \times n$

$$\mathbf{Y} = \begin{bmatrix} y_1 & & 0 \\ & \ddots & \\ 0 & & y_n \end{bmatrix}$$

$$\max_{\boldsymbol{\alpha} \geq 0} \boldsymbol{\alpha}^T \mathbf{1} - \frac{\boldsymbol{\alpha}^T (\mathbf{Y}^T \mathbf{X}^T \mathbf{X} \mathbf{Y}) \boldsymbol{\alpha}}{2}$$

Support Vectors

$$\boldsymbol{\alpha}^* = \begin{bmatrix} \alpha_1^* \\ \vdots \\ \alpha_n^* \end{bmatrix} \quad \mathbf{w}^* = \sum_{i=1}^n \alpha_i^* \mathbf{x}_i y_i$$

Support Vectors

$$\boldsymbol{\alpha}^* = \begin{bmatrix} \alpha_1^* \\ \vdots \\ \alpha_n^* \end{bmatrix}$$

$$\mathbf{w}^* = \sum_{i=1}^n \alpha_i^* \mathbf{x}_i y_i$$

Definition: A support vector is a point for which $\alpha_i^* > 0$

Support Vectors

$$\boldsymbol{\alpha}^* = \begin{bmatrix} \alpha_1^* \\ \vdots \\ \alpha_n^* \end{bmatrix} \quad \mathbf{w}^* = \sum_{i=1}^n \alpha_i^* \mathbf{x}_i y_i$$

Definition: A support vector is a point for which $\alpha_i^* > 0$

$$\alpha_i^* g_i(\mathbf{w}^*) = 0 \implies \alpha_i^* \left[1 - \left((\mathbf{w}^*)^T \mathbf{x}_i \right) y_i \right] = 0$$

Complementary Slackness

Support Vectors

$$\boldsymbol{\alpha}^* = \begin{bmatrix} \alpha_1^* \\ \vdots \\ \alpha_n^* \end{bmatrix} \quad \mathbf{w}^* = \sum_{i=1}^n \alpha_i^* \mathbf{x}_i y_i$$

$$\alpha_i^* g_i(\mathbf{w}^*) = 0 \implies \alpha_i^* \left[1 - \left((\mathbf{w}^*)^T \mathbf{x}_i \right) y_i \right] = 0$$

Complementary Slackness

Definition: A support vector is a point for which $\alpha_i^* > 0$

Every support vector lies on one of the two supporting hyperplanes $(\mathbf{w}^*)^T \mathbf{x} = \pm 1$

Support Vectors

$$\boldsymbol{\alpha}^* = \begin{bmatrix} \alpha_1^* \\ \vdots \\ \alpha_n^* \end{bmatrix} \quad \mathbf{w}^* = \sum_{i=1}^n \alpha_i^* \mathbf{x}_i y_i$$

$$\alpha_i^* g_i(\mathbf{w}^*) = 0 \implies \alpha_i^* \left[1 - \left((\mathbf{w}^*)^T \mathbf{x}_i \right) y_i \right] = 0$$


Complementary Slackness

Definition: A support vector is a point for which $\alpha_i^* > 0$

Every support vector lies on one of the two supporting hyperplanes $(\mathbf{w}^*)^T \mathbf{x} = \pm 1$

Every point that is **not** on one of the two supporting hyperplanes has $\alpha_i^* = 0$.

Hard-Margin, Linear-SVM




$$\hat{y} = \begin{cases} 1, & (\mathbf{w}^*)^T \mathbf{x} \geq 0 \\ -1, & (\mathbf{w}^*)^T \mathbf{x} < 0 \end{cases}$$


SVM
Solver

$$\boldsymbol{\alpha}^* = \begin{bmatrix} \alpha_1^* \\ \vdots \\ \alpha_n^* \end{bmatrix}$$

$$\mathbf{w}^* = \sum_{i=1}^n \alpha_i^* \mathbf{x}_i y_i$$



Soft-Margin, SVM




Not linearly separable

(1) Structural \rightarrow Hard-margin, Kernel-SVM

(2) Statistical (outliers)

Soft-Margin, SVM




Not linearly separable

- (1) Structural \rightarrow Hard-margin, Kernel-SVM
- (2) Statistical (outliers)

- Ideally, we want $(\mathbf{w}^T \mathbf{x}_i) y_i \geq 1$.
- Not true for outliers.
- Use a non-negative bribe to push them

$$(\mathbf{w}^T \mathbf{x}_i) y_i + \xi_i \geq 1$$

Soft-Margin, SVM



Not linearly separable

- (1) Structural \rightarrow Hard-margin, Kernel-SVM
- (2) Statistical (outliers)

- Ideally, we want $(\mathbf{w}^T \mathbf{x}_i) y_i \geq 1$.
- Not true for outliers.
- Use a non-negative bribe to push them


$$(\mathbf{w}^T \mathbf{x}_i) y_i + \xi_i \geq 1$$

$$\min_{\mathbf{w}} \quad \frac{\|\mathbf{w}\|^2}{2} + C \cdot \sum_{i=1}^n \xi_i$$

sub. to

$$\begin{aligned} (\mathbf{w}^T \mathbf{x}_i) y_i + \xi_i &\geq 1, & 1 \leq i \leq n \\ \xi_i &\geq 0, & 1 \leq i \leq n \end{aligned}$$

Soft-Margin, SVM



$$\min_{\mathbf{w}} \frac{\|\mathbf{w}\|^2}{2} + C \cdot \sum_{i=1}^n \xi_i$$

sub. to


$$\begin{aligned} (\mathbf{w}^T \mathbf{x}_i) y_i + \xi_i &\geq 1, & 1 \leq i \leq n \\ \xi_i &\geq 0, & 1 \leq i \leq n \end{aligned}$$

Not linearly separable

- (1) Structural \rightarrow Hard-margin, Kernel-SVM
- (2) Statistical (outliers)

- Ideally, we want $(\mathbf{w}^T \mathbf{x}_i) y_i \geq 1$.
- Not true for outliers.
- Use a non-negative bribe to push them

$$(\mathbf{w}^T \mathbf{x}_i) y_i + \xi_i \geq 1$$



Dual for Soft-SVM

$$\max_{0 \leq \alpha \leq C} \alpha^T \mathbf{1} - \frac{\alpha^T (\mathbf{Y}^T \mathbf{X}^T \mathbf{X} \mathbf{Y}) \alpha}{2}$$

- Only change is the constraints
- Earlier there was only a lower-bound
- Now, we have an upper-bound
- These are called box-constraints

Soft-Margin, SVM: Hinge-loss formulation

$$\min_{\mathbf{w}} \quad \frac{||\mathbf{w}||^2}{2} + C \cdot \sum_{i=1}^n \xi_i$$

sub. to

$$(\mathbf{w}^T \mathbf{x}_i) y_i + \xi_i \geq 1, \quad 1 \leq i \leq n$$

$$\xi_i \geq 0, \quad 1 \leq i \leq n$$

\equiv

$$\min_{\mathbf{w}} \quad \frac{||\mathbf{w}||^2}{2} + C \cdot \sum_{i=1}^n \max(0, 1 - (\mathbf{w}^T \mathbf{x}_i) y_i)$$

Regularization

Hinge Loss

Soft-Margin, SVM: Hinge-loss formulation

$$\min_{\mathbf{w}} \quad \frac{\|\mathbf{w}\|^2}{2} + C \cdot \sum_{i=1}^n \max(0, 1 - (\mathbf{w}^T \mathbf{x}_i) y_i) \quad (1)$$

Terms in the objective:

(1) $\frac{\|\mathbf{w}\|^2}{2}$ controls the width of the margin
Smaller the value of $\|\mathbf{w}\|$, wider the margin





(2) $\sum_{i=1}^n \max(0, 1 - (\mathbf{w}^T \mathbf{x}_i) y_i)$ is the hinge-loss. Wider the margin, larger the loss.

Soft-Margin, SVM: Hinge-loss formulation

$$\min_{\mathbf{w}} \frac{\|\mathbf{w}\|^2}{2} + C \cdot \sum_{i=1}^n \max(0, 1 - (\mathbf{w}^T \mathbf{x}_i) y_i)$$

(1) (2)

- (1) and (2) work in opposite directions
- If $\|\mathbf{w}\|$ decreases, the margin becomes wider, which increases the hinge-loss.
- C controls the tradeoff between (1) and (2):
 - If C is small, we are fine with a wide margin.
 - If C is large, we prefer a narrow margin.
 - If $C \rightarrow \infty$, we do not tolerate bribery at all.



Closing Remarks

- Everything cannot remain a black-box.

Closing Remarks

- Everything cannot remain a black-box.
- Sometimes it is necessary to peep into the box.

Closing Remarks

- Everything cannot remain a black-box.
- Sometimes it is necessary to peep into the box.
- Sometimes it is necessary to build your own box.

Closing Remarks

- Everything cannot remain a black-box.
- Sometimes it is necessary to peep into the box.
- Sometimes it is necessary to build your own box.
- But we must know when to stop:

Closing Remarks

- Everything cannot remain a black-box.
- Sometimes it is necessary to peep into the box.
- Sometimes it is necessary to build your own box.
- But we must know when to stop:
 - Model

Closing Remarks

- Everything cannot remain a black-box.
- Sometimes it is necessary to peep into the box.
- Sometimes it is necessary to build your own box.
- But we must know when to stop:
 - Model
 - NumPy over sklearn

Closing Remarks

- Everything cannot remain a black-box.
- Sometimes it is necessary to peep into the box.
- Sometimes it is necessary to build your own box.
- But we must know when to stop:
 - Model
 - NumPy over sklearn
 - Native Python over NumPy

Closing Remarks

- Everything cannot remain a black-box.
- Sometimes it is necessary to peep into the box.
- Sometimes it is necessary to build your own box.
- But we must know when to stop:
 - Model
 - NumPy over sklearn
 - Native Python over NumPy
 - C over Python

Closing Remarks

- Everything cannot remain a black-box.
- Sometimes it is necessary to peep into the box.
- Sometimes it is necessary to build your own box.
- But we must know when to stop:
 - Model
 - NumPy over sklearn
 - Native Python over NumPy
 - C over Python
 - Assembly language over C

Closing Remarks

- Everything cannot remain a black-box.
- Sometimes it is necessary to peep into the box.
- Sometimes it is necessary to build your own box.
- But we must know when to stop:
 - Model
 - NumPy over sklearn
 - Native Python over NumPy
 - C over Python
 - Assembly language over C
 - Build a computer
 - ...